


Shared AI Review Lessons: Setup Guide

For: engineering teams that use Claude Code and Cursor on the same repositories.

Goal: Turn PR review findings into shared, versioned lesson files that both Claude Code and Cursor read. Everyone runs a self-review against them before opening a PR, so past findings aren't repeated. Each tool still does its own open review, so you keep two independent views.

1. How it works



- **Single source of truth.** Lessons live in the repo, not in anyone's personal AI memory.
- **Shared automatically.** Lessons merge through normal PRs, and everyone gets them with `git pull`, whichever tool they use.
- **Keeps learning.** Every review adds or updates lessons, and a periodic cleanup keeps the set small and relevant.
- **Independent reviews stay independent.** The self-review does an open review *before* it reads any lessons.

2. Repository structure

CLAUDE.md	# main agent instructions (Claude Code reads this)
AGENTS.md -> CLAUDE.md	# symlink, so Cursor reads the same instructions
docs/lessons/	
README.md	# lesson format + writing rules
2026-09-api-validate-request-body.md	
...	
.claude/skills/	# Claude Code reads natively; Cursor reads it for compatibility
record-lessons/SKILL.md	# review findings → lesson files
self-review/SKILL.md	# two-pass review before opening a PR
harvest-lessons/SKILL.md	# one-time bootstrap from past PRs, memories, rules

Why this works for both tools

- **Skills:** Cursor loads skills from `.claude/skills/` as well as its own folders, so one folder serves both tools.
- **Instructions:** Cursor reads `AGENTS.md` and Claude Code reads `CLAUDE.md`. The symlink makes them the same file.

Windows teammates: Git symlinks need `git config core.symlinks true` and Windows Developer Mode. If that isn't possible, make `AGENTS.md` a normal file containing only:

Read and follow `CLAUDE.md` in this folder. It is the source of truth for agent instructions.

3. Setup steps (one person does this once per repo)

```
git checkout -b chore/shared-ai-lessons
mkdir -p docs/lessons .claude/skills/record-lessons .claude/skills/self-review
.claude/skills/harvest-lessons
ln -s CLAUDE.md AGENTS.md           # skip on Windows; use the pointer file above
```

Then:

1. Create the files in sections 4 through 8 with the contents shown.
2. Add the section 5 block to your existing `CLAUDE.md`. Don't replace the file.
3. If your default branch isn't `main`, replace `main` in the skills with its name.
4. Optionally, add the drift check from section 9 to CI.
5. Open a PR, get it reviewed and merged.
6. Run the harvest from section 10 so the team doesn't start with an empty lessons folder.

4. docs/lessons/README.md

Lessons

Each file is one lesson learned from a code review. Claude Code and Cursor both read these during `self-review`. Lessons are reviewed in PRs like code.

File name

`YYYY-MM-<area>-<short-kebab-name>.md`, e.g. `2026-09-api-validate-request-body.md`

Format

```markdown

---

```
title: Validate request bodies at the API boundary
paths: ["src/api/**"] # globs this lesson applies to; ["**"] if global
tags: [validation, security]
severity: high # high | medium | low
source: "PR #42 review" # where it came from
found_by: teammate # claude | cursor | teammate | <name>
added: 2026-09-14
updated: 2026-09-14
occurrences: 1 # bump when the same finding shows up again
```

---

### ## Rule

One or two sentences, phrased as a general rule (not a one-off fix).

### ## Why

What went wrong, or could go wrong. Reference the real incident briefly.

### ## How to check

What a reviewer (human or AI) should look for in a diff to catch this.

```

Writing rules

- One rule per file. If two lessons say the same thing, merge them and bump `occurrences`.
- General, not specific: "validate request bodies", not "fix line 40 of users.ts".
- Describe code problems, never people. No names in Rule/Why.
- No secrets, credentials, customer data, or internal URLs.
- Delete lessons that no longer apply (code removed, rule enforced by a linter, etc.).

5. Add to CLAUDE.md

Add this block to your existing CLAUDE.md :

Review lessons (shared by Claude Code and Cursor)

- Team review lessons live in `docs/lessons/`, one file per lesson. See `docs/lessons/README.md`.
- Never save PR review findings to personal memory. Record them in `docs/lessons/` with the `record-lessons` skill so the whole team (and Cursor) can use them.
- Before opening a PR, run the `self-review` skill.
- When you learn something durable about this codebase from a review, propose a lesson for it.

6. `.claude/skills/record-lessons/SKILL.md`

```
---
name: record-lessons
description: Turn PR review findings into shared lesson files in docs/lessons/. Use after a PR
review is received or given, when the user says "record lessons", "save this finding", "remember
this for next time", or pastes review comments.
---

# Record lessons from a review

1. Get the findings. Use, in order of preference:
  - review comments the user pasted;
  - a PR number/URL: fetch its review comments and review summaries (e.g. `gh pr view <n> --
comments`
    and `gh api repos/{owner}/{repo}/pulls/<n>/comments`);
  - findings from the current conversation.
2. Filter. Keep only findings that were valid and would apply to future code. Skip pure style
nits already enforced by linters, one-off typos, and anything the author rejected with a good
reason.
3. Generalize. Rewrite each kept finding as a general rule using the format in
`docs/lessons/README.md`.
  Set `paths` to the narrowest globs that cover where this could recur.
4. Deduplicate. Read the existing lessons in `docs/lessons/`. If one already covers the rule,
update it instead (add the new source, bump `occurrences`, refresh `updated`, sharpen "How to
check").
5. Write new files as `docs/lessons/YYYY-MM-<area>-<short-name>.md`. Set `found_by` to who
caught it
  (`claude`, `cursor`, `teammate`, or a name).
6. Check for sensitive content. No names in Rule/Why, no secrets, no customer data.
7. Report a short list: created / updated / skipped (with reason). Suggest committing the
lessons
  in the current PR or a small separate `chore: lessons` PR.
```

7. `.claude/skills/self-review/SKILL.md`

```
---
name: self-review
description: Review the current branch before opening a PR - an independent bug review first, then
a check against the team's lessons in docs/lessons/. Use when the user says "self review", "review
my branch", "check before PR", or is about to open a PR.
---

# Self-review before a PR

Do the two passes in order. Do NOT read docs/lessons/ until pass 1 is finished. This keeps the
open review independent instead of turning it into a checklist.

## Setup
- Base branch: `main` (change if the repo uses another).
- `git fetch origin main` then diff with `git diff origin/main...HEAD` and list files with
  `git diff --name-only origin/main...HEAD`. Include uncommitted changes (`git diff`) if any.

## Pass 1 - Open review (no lessons yet)
Read the diff and enough surrounding code to understand it. Look for:
- correctness bugs, logic errors, off-by-one, wrong conditions
- unhandled errors, null/undefined, edge cases, empty inputs
- security: auth/authz, injection, unvalidated input, secrets, unsafe data exposure
- concurrency, race conditions, transactions, idempotency
- performance traps (N+1 queries, unbounded loops/queries)
- missing or weak tests for the changed behavior
Record findings with file:line, why it is a problem, and a suggested fix.

## Pass 2 - Lessons check
1. List `docs/lessons/*.md` (skip README.md). Read the frontmatter of each.
2. Select lessons whose `paths` match any changed file, or whose `tags` clearly relate to the
change.
   Always include `severity: high` lessons with `paths: ["**"]`.
3. For each selected lesson, apply its "How to check" to the diff.
4. Record violations with file:line and the lesson file name.

## Report
- Pass 1 - Independent findings (most severe first)
- Pass 2 - Lesson violations (cite the lesson file)
- Lessons checked: count, and which were relevant
- Candidate new lessons: pass-1 findings that look like recurring patterns. Offer to record
them
  with the `record-lessons` skill.
Do not fix anything unless the user asks.
```

8. `.claude/skills/harvest-lessons/SKILL.md` (bootstrap, so you don't start empty)

```
---
name: harvest-lessons
description: Bootstrap docs/lessons/ from existing knowledge - past PR review comments, exported AI
memories, and existing rule files. Use when the user says "harvest lessons", "bootstrap lessons",
or pastes exported memories/rules to convert.
---

# Harvest lessons

Sources (use whichever the user provides or asks for):

## A. Past PR reviews in this repo
1. List recently merged PRs, e.g. `gh pr list --state merged --limit 100 --json
number,title,author`
   (ask the user how far back; default the last 100 merged PRs).
2. For each, fetch review comments and review bodies:
   `gh api repos/{owner}/{repo}/pulls/<n>/comments` and `gh api
repos/{owner}/{repo}/pulls/<n>/reviews`.
3. Ignore bot noise, approvals with no content, and resolved-as-wontfix threads.
4. Group comments that make the same point. A point raised in 2+ PRs is a strong lesson.
   A single high-severity finding (security, data loss) also qualifies.

## B. Exported memories / notes (Claude, Cursor, personal notes)
The user pastes text or points at a file. Extract only items about THIS codebase's
review findings, conventions, or pitfalls. Ignore personal preferences and anything about people.

## C. Existing rule files
Read `.cursor/rules/*`, `.cursorrules`, `CONTRIBUTING.md`, style guides, and any team review
checklists.
Convert review-relevant rules that aren't already enforced by linters.

## Then
1. Draft lessons using the format in `docs/lessons/README.md` (`source` = PR numbers or "harvest:
<source>").
2. Deduplicate against existing lessons and against each other.
3. Before writing, show the user a table: proposed title | severity | paths | sources |
occurrences.
   Let them drop or edit rows.
4. Write the approved lessons and suggest one `chore: harvest lessons` PR so the team can review
them.
```

9. Optional: drift check in CI

This makes sure `AGENTS.md` stays a symlink to `CLAUDE.md`:

```
[ "$(readlink AGENTS.md)" = "CLAUDE.md" ] || { echo "AGENTS.md must be a symlink to CLAUDE.md";
exit 1; }
```

If you used the pointer-file option for Windows, check that `AGENTS.md` mentions `CLAUDE.md` instead:

```
grep -q "CLAUDE.md" AGENTS.md || { echo "AGENTS.md must point to CLAUDE.md"; exit 1; }
```

10. Harvesting: don't start cold

Do this once after the setup PR merges. Each step ends in a PR, so the team reviews the lessons before they take effect.

10a. From past PR reviews (one person)

Requires the `gh` CLI, logged in with access to the repo. In Claude Code or Cursor, ask:

Use the harvest-lessons skill on the last 100 merged PRs in this repo.

Review the proposed table, approve it, and open the `chore: harvest lessons` PR.

10b. From your own Claude memories (every teammate who saved review findings)

AI memories are personal and live on each person's computer or account. Nobody else can harvest them for you.

1. **Claude Code:** in the repo, run `/memory` to find and open your memory files. The automatic project memory is stored under `~/.claude/projects/`. Then ask:

Go through your memory for this project. Use the harvest-lessons skill to turn every PR review finding and codebase pitfall into lesson files in `docs/lessons/`. Show me the table first. After I approve and the lessons are written, remove those items from your memory.

1. **Claude app (claude.ai):** if you saved findings there, open Settings → memory, copy the relevant entries into a text file, and give them to the harvest skill (step 10c's prompt works).

10c. From Cursor memories and rules (every Cursor user)

1. Copy any review-related entries from Cursor's settings (Rules / Memories / User Rules) into a text file.
2. Ask, in either tool:

Use the harvest-lessons skill on this exported text: `<paste or path>`. Only keep items about this codebase.

1. Also check whether the repo has `.cursor/rules/` or `.cursorrules`. The skill can harvest those directly.

10d. Submit

Each person opens a `chore: harvest lessons (<name>)` PR. Reviewers merge duplicates. The skill dedupes, but check anyway.

11. Daily workflow

When	You say (in Claude Code or Cursor)	What happens
You get a PR review	"Record lessons from PR #123"	New or updated lesson files to commit
You review a teammate's PR	"Record lessons from my review on PR #456"	Same; commit in a small <code>chore: lessons</code> PR
Before opening a PR	"Self review"	Pass 1 open review, pass 2 lessons check
Monthly	"Review docs/lessons: merge duplicates, delete stale ones"	Smaller, sharper lesson set

Tip: Cursor may pick up skills automatically less reliably than Claude Code. Asking by name ("use the self-review skill") makes it dependable.

12. Keeping Claude and Cursor reviews independent

The shared lessons make both tools catch the same *known* issues, which is intended. To keep their *independent* findings different:

- Two-pass self-review.** Pass 1 runs before any lessons are read. This is built into the skill.
- Different model families.** Use a non-Claude model in Cursor (for example GPT or Gemini). If Cursor runs a Claude model, the two reviews end up much more alike.
- Don't share results early.** Run each review in a fresh chat, and compare only after both are done.
- Record who caught it.** The `found_by` field shows which tool or person contributes which lessons. Lessons from one tool's catches teach the other.

13. Rules of thumb

- Repo vs. personal memory:** Facts about the codebase go in the repo. Personal preferences (tone, formatting) stay in personal memory.
- Lessons are public to the repo:** Anyone with repo access can read them, and they stay in git history. No secrets, customer data or names.
- Don't duplicate linters:** If a linter or type check can enforce a lesson, add the lint rule and delete the lesson.

- **Keep the set small:** Fewer, sharper lessons beat a long list nobody, human or AI, reads carefully.